

Adversarial Reachability for Program-level Security Analysis

Anonymous authors

Abstract. Many program analysis tools and techniques have been developed to assess program vulnerability. Yet, they are based on the standard concept of reachability and represent an attacker able to craft smart *legitimate* input, while in practice attackers can be much more powerful, using for instance micro-architectural exploits or fault injection methods. We introduce *adversarial reachability*, a framework allowing to reason about such *advanced attackers* and check whether a system is vulnerable or immune to a particular attacker. As equipping the attacker with new capacities significantly increases the state space of the program under analysis, we present a new symbolic exploration algorithm, namely *adversarial symbolic execution*, injecting faults in a *forkless* manner to prevent path explosion, together with optimizations dedicated to reduce the number of injections to consider while keeping the same attacker power. Experiments on representative benchmarks from fault injection show that our method significantly reduces the number of adversarial paths to explore, allowing to scale up to 10 faults where prior work timeout for 3 faults. In addition, we analyze the well-tested WooKey's bootloader, and demonstrate the ability of our analysis to find attacks and evaluate countermeasures in real-life security scenarios. We were especially able to find a new attack on an incomplete patch.

Keywords: Program analysis · Attacker model · Fault injection · Symbolic execution

1 Introduction

Context Major works have delved into program analysis over the last decades, leveraging techniques such as symbolic execution [18, 24, 53], static analysis [43], abstract interpretation [30] or bounded model checking [29], to hunt for software vulnerabilities and bugs in programs, or to prove their absence [35, 60], leading to industrial adoption in some leading companies [6, 18, 43, 60, 66]. As bugs are an attack entry point, removing them is a first step towards better software security.

Problem Yet, stepping back from these successes, it appears that all these methods consider a rather weak threat model, where the attacker can only craft smart “inputs of death” through legitimate input sources of the program, exploiting corner cases in the code itself. Tools only looking for bugs and software vulnerabilities may deem a program secure while the bar remains quite low for an *advanced attacker*, able for example to take advantage of attack vectors such as (physical) hardware fault injections [58], micro-architectural attacks [61, 70],

37 software-based hardware attacks [55,69,86] like Rowhammer [70], or any com-
 38 bination of vectors [63]. While previously limited to high-security devices and
 39 systems such as smart cards and cryptography modules [13,16], these fault-based
 40 attacks can now target a wider spectrum of systems, such as bootloaders [57],
 41 firmware update modules [19], security enclaves [69], etc. The reasoning behind
 42 automated software-implemented fault injection also applies to Man-At-The-End
 43 attacks [3] and is similar to the (manual) reasoning performed in control-flow
 44 integrity to evaluate countermeasures [1,21].

45 **Goal & Challenges** *Our goal is to devise a technique to automatically and*
 46 *efficiently reason about the impact of an advanced attacker onto a program se-*
 47 *curity properties*, where the standard reachability framework only supports an
 48 attacker crafting smart legitimate inputs. The first challenge is to provide a formal
 49 framework to study what an advanced attacker can do to attack a program
 50 under study. Interestingly, while such frameworks are routinely used in crypt-
 51 ographic protocol verification [7,26], none has been studied for program-level
 52 analysis. The second challenge is to design an efficient algorithm to assess the
 53 vulnerability of a program to a given attacker model, while adding capabilities
 54 to the attacker naturally give rise to a significant path explosion – especially in
 55 the case of multiple fault analysis.

56 The rare prior work in the field, mostly focused on encompassing physi-
 57 cal fault injections for high-security devices, rely mostly on *mutant genera-*
 58 *tion* [25,28,49,50,79] or *forking analysis* [15,20,63,76], yielding scalability issues.
 59 Moreover, most of them are limited to a few predefined fault models and do not
 60 propose any formalization of the underlying problem.

61 **Proposal** We propose *adversarial reachability*, a formalism extending standard
 62 reachability to reason about a program execution in the presence of an advanced
 63 attacker, and we build a new algorithm based on symbolic techniques, named
 64 *adversarial symbolic execution*, to address the adversarial reachability problem
 65 from the bug finding point of view (bounded verification). Our algorithm pre-
 66 vents path explosion thanks to a new *forkless* encoding of faults. We show it
 67 correct and k-complete with respect to adversarial reachability. To improve the
 68 performance further, we design two new optimizations to reduce the number of
 69 injected faults: Early Detection of fault Saturation and Injection On Demand.

70 **Contributions** As a summary, we claim the following novelties:

- 71 – We formalize the adversarial reachability problem (Section 4), extending
 72 standard reachability to take into account an advanced attacker, together
 73 with the associated correctness and completeness definitions;
- 74 – We describe a new symbolic exploration method (Section 5), adversarial sym-
 75 bolic execution, to answer adversarial reachability, featuring a novel forkless
 76 fault encoding to prevent path explosion and two optimization strategies to
 77 reduce fault injection. We establish their correctness and completeness;
- 78 – We propose an implementation of our techniques for binary-level analysis
 79 (Section 6), on top of the BINSEC framework [38]. We systematically evalu-
 80 ate its performances against prior work (Section 7), using a standard SWiFI
 81 benchmark from physical fault attacks and smart cards. Experiments show

82 a very significant performance gain against prior approaches, for example up
 83 to x10 and x215 times on average for 1 and 2 faults respectively – with a
 84 similar reduction in the number of adversarial paths. At most, we are x224
 85 times faster for 1 fault and x6121 for 2. Moreover, our approach scales up to
 86 10 faults whereas the state-of-the-art starts to timeout for 3 faults.
 87 – We finally perform a security analysis of the WooKey bootloader¹ (Section
 88 8), a very well tested real-life security focused program. We were able to
 89 replay known attacks, find new ones and evaluate the adequacy of some
 90 of the countermeasures. Especially, we found a new attack on a recently
 91 proposed patch [63], and proposed a patch to the developers.

92 This work is a first step in designing efficient program analysis techniques able to
 93 take into account advanced attackers. The approach is generic enough to accom-
 94 modate many common fault models, including the bit flip from RowHammer,
 95 test inversion or arbitrary data modification; still, instruction skips or modifica-
 96 tions are currently out of reach. Also, while we investigate the bug finding side
 97 of the problem (underapproximation), the verification side (overapproximation)
 98 is interesting as well. These are exciting directions for future research.

99 *Our dataset and benchmark infrastructure will be made available through ar-*
 100 *tifact for reproducibility purpose, and the code will be open-sourced.*

101 2 Motivation

102 We start by motivating the need for adversarial reachability, first with a descrip-
 103 tion of several realistic attack scenarios involving advanced attackers (Section
 104 2.1), second with a small example showing the need for dedicated analysis (Sec-
 105 tion 2.2).

106 2.1 Fault injection across security fields

107 We describe hereafter several real software-level security scenarios where the
 108 attacker goes beyond crafting legitimate input to abuse the system under at-
 109 tack. Interestingly, while these scenarios were historically focused on hardware-
 110 hardened high-security systems (such as smart cards) and associated with com-
 111 plex physical attack means, many recent scenarios do involve software-only at-
 112 tacks on standard systems, with targets encompassing cryptographic libraries,
 113 bootloaders, firmware updaters, security enclaves, etc.

114 **Hardware fault injection attacks** [58] cause erroneous computations by dis-
 115 turbing the signal propagation in the chip with physical means such as elec-
 116 tromagnetic pulses [39], laser beams [4, 85], or power [19] and clock glitches.
 117 The associated fault models include bit-, byte- or word- set and reset, bit-flips,

¹ WooKey [14, 88] is a secure USB mass storage device developed by the French Na-
 tional Security Agency, and has recently served as a recent challenge among French
 security evaluators.

118 instructions corruption and instruction skips. State-of-the-art attacks involve
119 multiple fault injections [59], as expected by the high level of attack potential in
120 Common Criteria vulnerability analysis.

121 **Software-implemented hardware attacks** push the hardware into unstable
122 states using software controlled mechanisms, like delays in memory buses in-
123 ducing bit-flips in data fetched from memory [55] or CPU voltage and frequency
124 manipulations yielding bit-flips in the processor [69, 86]. The notorious *Rowham-*
125 *mer* attack [70] abuses memory accesses to induce bit-flips in flash memory.

126 **Micro-architectural attacks** use micro-architectural behaviors in unexpected
127 ways. For example: the Spectre attack (version v1) [62] exploits branch predictors
128 in speculative executions, which can be seen as a test inversion followed by a
129 rollback; race attacks [54] corrupt data of other running processes and can be
130 seen as arbitrary data faults.

131 **Man-At-The-End (MATE) attacks** considers an attacker having full observ-
132 ability and control over a software code and its execution [3], with the goal to
133 steal sensitive data or code (reverse engineering attacks). The associated attacker
134 model is hence very powerful, with capabilities such as halting and modifying
135 data and code at any point of the execution.

136 **CFI reasoning** In order to assess the power of Control-Flow Integrity (CFI)
137 mechanisms, researchers [1, 21] define hypothetical attackers by their capabilities,
138 such as “write anything anywhere” or “write anything somewhere”, and manually
139 prove that their countermeasure is indeed able to thwart such an opponent.
140 While not *per se* an applicative security scenario, the techniques developed in
141 this paper could help automate such essential reasoning.

142 2.2 Motivating example

143 The motivating example in Figure 1 is a simple unrolled program inspired by
144 the VerifyPIN benchmark [42], from the domain of hardware fault injection and
145 smart cards. The user PIN digits *u1* to *u4* are checked against the reference digits
146 *ref1* to *ref4*, using the accumulator *res*. The attacker seeks to be authenticated
147 (validate the assert l.16) without knowing the right digits (l.14).

148 Here, the attacker indeed cannot succeed by only crafting legitimate inputs.
149 However, an advanced attacker can leverage more powerful attack vectors to
150 inject faults into the program in order to succeed. For instance, corrupting
151 *g_authenticated* to *true* at l.10 achieves the attacker goal. It could be obtained
152 for example through a physical- or Rowhammer- attack.

153 **Program analysis** As expected, standard symbolic execution tools such as
154 Klee [22], angr [84] or BINSEC [38] do not report any violation here, as they
155 consider the simplest possible attacker. We can try to use software-implemented
156 fault injection (SWiFI) techniques [15, 20, 63, 76] (detailed in Section 3.1) from
157 the high-security system evaluation. Yet, the standard *forking* approach does
158 not scale with multiple faults: here, 166 paths are explored in 0.6 seconds for 1
159 fault, 2994 paths in 11 seconds for 2 faults, and it keeps on adding a factor x10 in

```

1 bool g_authenticated;
2 int u1, u2, u3, u4, ref1, ref2, ref3, ref4;
3
4 void verifyPIN () {
5     int res = 1;
6     res = res * (u1 == ref1);
7     res = res * (u2 == ref2);
8     res = res * (u3 == ref3);
9     res = res * (u4 == ref4);
10    g_authenticated = res;
11 }
12
13 void main(int argc, char const *argv[]) {
14     assert(u1!=ref1 || u2!=ref2 || u3!=ref3 || u4!=ref4);
15     verifyPIN();
16     assert(g_authenticated == true); /* Security oracle */
17 }

```

Fig. 1: Motivating example, inspired by VerifyPIN [42]

160 explored paths and analysis time for each extra fault, until the analysis timeouts
161 (12 hours) above 4 faults. On the contrary, our *forkless* algorithm presented
162 in Section 5 simulates fault injection without creating new paths and, in this
163 example, shows a constant runtime as the number of faults increases from 1 to
164 10 – we explore 9 paths in 0.2 seconds in all cases.

165 3 Background

166 We provide in this section background information on software-implemented
167 fault injection, standard reachability and symbolic execution.

168 3.1 Software-implemented fault injection (SWiFI)

169 Software-implemented fault injection (SWiFI) tools [15, 20, 25, 28, 49, 50, 63, 68,
170 76, 79] have been developed in the community of high-secure systems to ease
171 hardware fault injection campaigns, which are time consuming and require spe-
172 cial equipment. SWiFI evaluates a program with the transformations induced
173 by the effects of hardware faults, in order to find interesting attack paths. We
174 distinguish two main SWiFI techniques.

175 First, the *Mutant generation* approach [25, 28, 49, 50, 79] consists in analysing
176 slightly modified versions of the program (named mutants), each on them embed-
177 ding a different faulty instruction. For example, a few mutants of the motivating
178 example are presented in Appendix A, Figure 7. Each mutant is then analyzed
179 on its own, typically with symbolic execution. The main limitation of mutant
180 generation is the explosion of mutants, in particular for multiple faults. Also,

181 as the different mutants differ only slightly, analyzing each of them separately
 182 wastes lots of time repeating similar reasoning.




Fig. 2: Forking technique transformation in pseudo-code

183 Second, the *forking approach* [15, 20, 63, 76] consists in instrumenting the
 184 analysis (or the code, via instrumentation) to add all possible faults as forking
 185 points (branches) controlled by boolean values indicating whether a particular
 186 fault will be taken or not, plus constraints on the maximal number of faults
 187 allowed. A forking data fault is illustrated in Figure 2. An instrumented version
 188 of the motivation example is presented in Appendix A, Figure 8. A standard
 189 program analysis technique is then launched – typically symbolic execution or
 190 bounded model checking. Compared with mutant generation, this method allows
 191 to share the analysis between the different possible faults. Still, the number of
 192 paths explodes with the number of possible faults (forking points).

193 **Scalability issues** These two approaches yield an explosion of the whole search
 194 space w.r.t. the number of fault injection points in the program: the mutant
 195 approach leads to consider up to C_k^n (k among n)¹. mutants for a program
 196 under analysis with n possible fault locations and k faults, while the forking
 197 approach yields up to C_k^n paths to analyzed for a single original program path
 198 with n possible fault locations and k faults.

199 *In the following, we will consider the forking approach as the baseline – please*
 200 *keep in mind that the mutant approach scales worstly.*

201 **Fault models** Supported fault models vary for each tool, but they are usually
 202 adapted from hardware fault models [47, 82]. The most common fault models
 203 are (1) *data faults* such as arbitrary data faults, set and reset of bytes, words or
 204 variables, bit-flips; as well as (2) *instruction corruptions* such instruction skips
 205 and test inversions. Most tools are limited to one (sometimes two) hard-coded
 206 fault models. Only few SWiFI tools can handle multiple faults [63, 68, 76, 87] –
 207 still with scalability issues.

208 3.2 Standard reachability formalization

209 Considering a program P , we denote S the set of all possible states of P . A state
 210 is composed of the code memory, the data memory (i.e. the stack and heap),
 211 the state of registers and the location of the next instruction to execute. The

¹ Remind that $C_k^n = \binom{n}{k} = \frac{n!}{k!(n-k)!}$

212 set of input states of a program P is noted $S_0 \subset S$. The set of transitions (or
 213 instructions) of the program is denoted T . The execution of an instruction t is
 214 represented by a one-step transition relation $\rightarrow_t \in S \times S$. We denote $s \rightarrow s'$ when
 215 $s \rightarrow_t s'$ for some $t \in T$. We extend the transition relation over any finite path
 216 $\pi \in T^*$ through composition. The transitive reflexive closure of \rightarrow is noted \rightarrow^* .
 217 Finally, we use $S \rightarrow s'$ as a shortcut for $\exists s \in S. s \rightarrow s'$, and $\rightarrow_{\leq k}$ for reachability
 218 in at most k steps.

219 We consider in the rest of the paper the case of *location reachability*: given
 220 a location l (instruction or code address) of the program under analysis, the
 221 question is whether we can reach any state s at location l . More formally, L is
 222 the finite set of locations of P , and we consider a mapping $loc : S \mapsto L$ from
 223 states to locations. For example, loc may return the program counter value. We
 224 write $S \rightarrow^* l$ as a shortcut for $\exists s' \in S. S \rightarrow^* s' \wedge loc(s') = l$.

225 **Definition 1 (Standard reachability).** *A location l is reachable in a program*
 226 *P if $S_0 \rightarrow^* l$.*

227 We now define correctness and completeness for a program analyzer.

228 **Definition 2 (Correctness, completeness).** *Let $\mathcal{V} : (P, l) \mapsto \{1, 0\}$ be a*
 229 *verifier taking as input a program P and a target location l .*

- 230 – \mathcal{V} is correct when for all P, l , if $\mathcal{V}(P, l) = 1$ then l is reachable in P ;
- 231 – \mathcal{V} is complete when for all P, l , if l is reachable then $\mathcal{V}(P, l) = 1$;
- 232 – if \mathcal{V} also takes an integer bound n as input, \mathcal{V} is k -complete when for all
 233 bound n and P, l , if l is reachable in at most n steps then $\mathcal{V}(P, l, n) = 1$.

234 We want to stress out that while location reachability can be seen as a basic
 235 case, we consider it sufficient here for two reasons: first, it keeps the formalism
 236 light while still straightforward to generalize to stronger reachability properties
 237 (e.g., local predicates of the form (l, φ) , sets of finite traces, etc.); second, it
 238 is already rather powerful on its own, as we can still instrument the code to
 239 reduce some stronger forms of reachability to it (e.g., adding local assertions or
 240 monitors).

241 3.3 Symbolic execution

242 Symbolic execution (SE) [23, 24, 52, 83] is a symbolic exploration technique for
 243 standard reachability. Algorithm 1 gives a high-level view of a typical SE al-
 244 gorithm, adapted for location reachability². The analysis follows each possible
 245 path π of a program up to a depth bound k . If π reaches the target, then we
 246 check whether π is indeed feasible by computing its *path predicate* Φ – a logical
 247 formula representing the path constraints over the input variables along π , and
 248 sending it to a SMT solver [12], that will try to answer whether the formula
 249 is satisfiable or not, and provide a model for free variables (e.g. inputs) if it is
 250 (omitted here for simplicity). SE is *correct* for location reachability, and even
 251 *k-complete* if we assume a perfect encoding of path predicates.

Algorithm 1: Standard symbolic execution algorithm, taken from [48]

Input: a program P , a bound k , a target location l
Output: Boolean value indicating whether l can be reached within k steps.

```

1 for path  $\pi$  in GetPaths( $k$ ) do
2   if  $\pi$  reaches  $l$  then
3      $\Phi := \text{GetPredicate}(\pi)$ 
4     if  $\Phi$  is satisfiable then
5       return true
6     end
7   end
8 end
9 return false

```

Algorithm 2: Assignment evaluation in SE

Input: path predicate Φ , assignment instruction $x := \text{expr}$
Output: Updated Φ

```

1 Function eval_assign( $\Phi, x, \text{expr}$ ) is
2   return  $\Phi \wedge (x \triangleq \text{expr})$ 
3 end

```

252 In this paper, we will focus on the evaluation of assignments and conditional
253 jumps for SE, detailed in Algorithms 2 and 3 respectively, as this is where our ad-
254 versarial symbolic execution will mainly differ from the standard one. It requires
255 going slightly deeper into details. In practice, the program paths are explored
256 incrementally. A worklist WL records all pending paths together with their as-
257 sociated path predicate and their next instruction to explore. On conditional
258 branches, the symbolic path is split in two (one for each branch, updating the
259 path constraint accordingly), and each new prefix is added to the worklist (Al-
260 gorithm 3). Assignments are dealt with straightforwardly, simply adding a new
261 logical variable definition to the path predicate ³ (notation: $x \triangleq y$).

262 4 Adversarial reachability

263 In this section, we detail the advanced attacker model we consider and define
264 the adversarial reachability problem. Especially, *advanced attackers can do more*
265 *than carefully crafting legitimate inputs to trigger vulnerabilities in a software.*

² More complex properties can be verified with the same principles, such as local predicate reachability, trace properties or hyper-properties [36].

³ Actually, a symbolic state usually comprises the path predicate itself plus a mapping from program variable names to logical variable names, and assignments involve both creating new logical names and updating the mapping. We abstract away from these details.

Algorithm 3: Conditional jump evaluation in SE

Input: path predicate Φ , conditional jump instruction *if cdt then l_t else l_e*
Data: a worklist WL containing the pending path prefixes to explore – list of pairs (path predicate, next location)
Output: WL updated in place

```

1 Function eval_conditional_jump( $\Phi$ ,  $cdt$ ,  $l_t$ ,  $l_e$ ) is
2   if  $\Phi \wedge cdt$  is satisfiable then
3     | Add ( $\Phi \wedge cdt$ ,  $l_t$ ) to  $WL$ 
4   end
5   if  $\Phi \wedge (\neg cdt)$  is satisfiable then
6     | Add ( $\Phi \wedge \neg cdt$ ,  $l_e$ ) to  $WL$ 
7   end
8 end

```

266 They can use a wide variety of attack vectors (e.g. hardware fault injection at-
 267 tacks, software-implemented hardware attacks, micro-architectural attacks, soft-
 268 ware attacks, etc), in any combination, and multiple times. We suppose attack
 269 vectors prerequisites have been met, and only consider the impact of the faults
 270 on the program under attack.

271 Our *attacker model* has three components: (1) a set of attacker actions, equiv-
 272 alent to fault models; (2) the maximum number of actions the attacker can
 273 perform; and (3) a goal, expressed here as a location reachability query.

274 Formally, given a program P with set of states S , set of transitions T and set
 275 of locations L , we extend the transition model described in Section 3.2 to include
 276 an adversarial transition $\leadsto_A \in S \times S$ related to an attacker A , i.e. $T_A = T \cup \leadsto_A$.
 277 To specify practical fault models, restrictions are applied onto \leadsto_A , limiting
 278 what part of the state can be modified and how. For instance, when considering
 279 arbitrary data faults, only the data memory and the register values could be
 280 modified. Then, the transition relation of P under attacker A is denoted as
 281 $\mapsto_A = \rightarrow \cup \leadsto_A = (\cup_{t \in T} t) \cup \leadsto_A$. We extend the notations from Section 3.2 to
 282 the relation \mapsto_A . Especially, $S \mapsto_A^* s'$ means $\exists s \in S. s \mapsto_A^* s'$, the adversarial
 283 transition relation up to k is denoted $\mapsto_{A, \leq k}$.

284 Still, we need to take into account the maximum number of faults the at-
 285 tacker can perform along an execution. Given a path π over T_A^* , π is said to be
 286 *legit* if it does not contain \leadsto_A , and *faulty* otherwise. The number of occurrences
 287 of transition \leadsto_A in π is its *number of faults*. Given a bound m_A on the fault
 288 capability of A , we define $\mapsto_{(A, m_A)}^*$ by limiting the adversarial reachability rela-
 289 tion to paths π with less than m_A faults. We consider m_A to be $+\infty$ in case the
 290 attacker has no such limitation. For the sake of simplicity, in the following, we
 291 will consider m_A as an implicit parameter of A , and simply write \mapsto_A^* instead of
 292 $\mapsto_{(A, m_A)}^*$.

293 **Definition 3 (Adversarial reachability).** *Given an attacker A with a m_A*
 294 *faults budget and a program P , a location $l \in L$ is adversarially reachable if*
 295 *$S_0 \mapsto_A^* s' \wedge \text{loc}(s') = l$ for some $s' \in S$.*

296 In the following, adversarial reachability of location l from a set of states S
 297 will be denoted $S_0 \mapsto_A^* l$.

298 **Proposition 1.** *Standard reachability implies adversarial reachability. The con-*
 299 *verse does not hold.*

300 *Proof.* Standard reachability can be viewed as adversarial reachability with an
 301 attacker able to perform 0 faults.

302 We redefine what it means for an analysis answering adversarial reachability
 303 to be correct, complete and k -complete.

304 **Definition 4.** *Let $\mathcal{V}_A : (P, A, l) \mapsto \{1, 0\}$ be a verifier taking as input a program*
 305 *P , an attacker A with m_A fault budget and a target location l .*
 306 *– \mathcal{V}_A is correct given A when for all P, l , if $\mathcal{V}_A(P, A, l) = 1$ then l is adver-*
 307 *sarially reachable in P for attacker A ;*
 308 *– \mathcal{V}_A is complete given A when for all P, l , if l is adversarially reachable for*
 309 *attacker A then $\mathcal{V}_A(P, A, l) = 1$;*
 310 *– if \mathcal{V}_A also takes an integer bound n as input, \mathcal{V}_A is k -complete given A when*
 311 *for all integer n and P, l , if l is adversarially reachable in at most n steps*
 312 *then $\mathcal{V}_A(P, A, l, n) = 1$.*

313 5 Forkless Adversarial Symbolic Execution (FASE)

314 In this section, we present our forkless algorithm for adversarial reachability. The
 315 analysis aims to find inputs and a fault sequence compatible with the considered
 316 attacker model and reaching the target location. Our primary goal is to deal
 317 with the potential path explosion induced by possible faults. Our design guiding
 318 principles are the following:

- 319 – First, prevent path explosion as much as possible with a forkless fault en-
 320 coding. Yet, this forkless encoding leads to logical formulas potentially more
 321 complex and harder to solve in practice;
- 322 – Second, reduce as much as possible the complexity of the created formulas,
 323 by avoiding the undue introduction of extra-faults along a path.

324 5.1 Modelling faults via forkless encoding

325 The forkless encoding aims to address the path explosion induced by the forking
 326 treatment of fault injection in prior works. It is designed mainly for data faults
 327 and consists of wrapping arithmetically an assignment right-end side, as shown
 328 in Figure 3 for an arbitrary data fault. The activation of this fault location is
 329 determined by the symbolic Boolean value *fault_here*, and the corrupted value
 330 of x is the fresh variable *fault_value*.

331 The point is to embed the fault injection as an expression inside the logical
 332 formula, without any explicit path forking at the analysis top-level, in order to
 333 let the analyzer reason about both legit executions and faulty executions at the
 334 same time – this is akin to path merging in some ways, except that we do it only
 335 for the treatment of fault injection (we could also see the approach as avoiding
 336 undue path splits).

337 Multiple forkless arbitrary data encodings are possible. We chose to use the
 338 *ite* expression operator, an inlined form of if-then-else at the expression level.
 339 We also tried encodings inspired from branchless programming idioms (e.g.:
 340 $(b) \cdot x + (1 - b) \cdot y$. for *ite*(*b*, *x*, *y*) with *b* a Boolean value) – in our experiments
 341 they worked as well as the *ite* operator. Other data fault models are supported,
 342 such as set, reset, bit-flips, etc. Test inversion is also supported by applying
 343 faults to the condition of conditional jumps. Tables 1 illustrates various forkless
 344 encodings. Note that the forkless encoding is not designed for instruction cor-
 345 ruptions or instruction skips, as these modifications either yield permanent code
 346 modification or span several instructions.

$x := \text{expr}$	$x := \text{ite } \textit{fault_here} ? \textit{fault_value} : \text{expr}$
--------------------	---

(a) Original statement (b) Forkless transformation for arbitrary data fault

Fig. 3: Forkless injection technique

Table 1: Forkless encodings for various fault models

Fault model	original instruction	Forkless encoding
Arbitrary data	$x := \text{expr}$	$x := \text{ite } \textit{fault_here} ? \textit{fault_value} : \text{expr}$
Variable reset	$x := \text{expr}$	$x := \text{ite } \textit{fault_here} ? 0x00000000 : \text{expr}$
Variable set	$x := \text{expr}$	$x := \text{ite } \textit{fault_here} ? 0xffffffff : \text{expr}$
Bit-flip	$x := \text{expr}$	$x := \text{ite } \textit{fault_here} ?$ $(\text{expr } \text{xor } 1 << \textit{fault_value}) : \text{expr}$
Test inversion	$\text{if } \textit{cdt} \text{ then goto 1}$ else goto 2	$\text{if } (\text{ite } \textit{fault_here} ? !\textit{cdt} : \textit{cdt})$ $\text{then goto 1 else goto 2}$

347 **Trade-off.** While these sorts of encoding indeed allow a significant path re-
 348 duction compared to forking approaches, the corresponding path predicates are
 349 more complicated than standard path predicates, as they involve lots of extra-
 350 symbolic variables for deciding whether the faults occur and emulating their
 351 effect. We show later in this section how to reduce these extra-variables.

352 5.2 Building adversarial path predicates

353 Adversarial symbolic execution requires modifications to Algorithms 2 and 3, as
 354 illustrated in Algorithms 4 and 5 respectively.

Algorithm 4: Forkless assignment evaluation

Input: path predicate Φ , assignment instruction $x := expr$

Output: Updated Φ

```

1 Function eval_assign( $\Phi, x, expr$ ) is
2   |    $\Phi', expr' := \text{FaultEncoding}(\Phi, expr)$ 
3   |   return  $\Phi' \wedge (x \triangleq expr')$ 
4 end
```

Algorithm 5: Forkless conditional jump evaluation

Input: path predicate Φ , conditional jump instruction *if* cdt l_t *else* l_e

Data: fault counter nb_f , maximal number of faults max_f , worklist WL

Output: WL updated in place

```

1 Function eval_conditional_jump( $\Phi, cdt, l_t, l_e$ ) is
2   |   if  $\Phi \wedge cdt \wedge (nb_f \leq max_f)$  is satisfiable then
3   |     |   Add ( $\Phi \wedge cdt, l_t$ ) to  $WL$ 
4   |   end
5   |   /* Idem for else branch ( $\neg cdt$ ) */
6 end
```

355 The assign evaluation process embeds a wrapper encoding the fault in a fork-
 356 less manner. Note that *FaultEncoding* involves the declaration of fresh symbolic
 357 variables for fault decisions and fault effects – hence the update of the path pred-
 358 icate Φ . Also, the fault counter nb_f is updated, and a new potentially faulted
 359 expression $expr'$ is computed.

360 Note that checking if the fault counter nb_f does not exceed the maximal
 361 number of faults max_f can be performed at different places. We found the best
 362 trade-off is to augment the conditional jump queries to check if we could explore
 363 each branch without exceeding max_f (see Algorithm 5), as checking at the end
 364 of a path often involves exploring many unfeasible faulty paths.

365 We refer to this set of modifications as *Forkless Adversarial Symbolic Execu-*
 366 *tion (FASE)*.

367 5.3 Algorithm properties

368 We now consider the properties of the FASE algorithm.

369 **Proposition 2.** *The FASE algorithm is correct and k -complete for adversarial*
 370 *reachability.*

371 *Sketch of proof.* If our algorithm finds an adversarial path reaching the target lo-
 372 cation l , by providing specific input values and a fault sequence, then an attacker
 373 executing the program with the provided inputs and performing the proposed
 374 faults will reach its goal. Our algorithm is based on symbolic execution with
 375 bounded path depth and explores all possible attack paths according to the
 376 considered attacker model, hence its k -completeness for adversarial reachability.

377 **Tightness of FASE.** Consider a single path with no branching instruction
 378 and an assert statement to be checked at the end, together with f possible fault
 379 locations and a maximum of m faults. Then the forking SE yields up to C_m^f paths
 380 to analyze, and as many queries to send to the solver. In the same scenario, FASE
 381 will analyze only the original program, and *send a single query to the solver.*

382 Still, the Forkless encoding increases query complexity, as shown in section
 383 7. We present in the remainder of this section two mitigation techniques.

384 5.4 Optimization via early detection of fault saturation (FASE-EDS)

Algorithm 6: FASE-EDS conditional jump evaluation

Input: path predicate Φ , conditional jump instruction *if* cdt *then* l_t *else* l_e
Data: fault counter nb_f , maximal number of faults max_f , worklist WL
Output: WL updated in place

```

1 Function eval_conditional_jump_EDS( $\Phi$ ,  $cdt$ ,  $l_t$ ,  $l_e$ ) is
2   if  $\Phi \wedge cdt \wedge (nb_f < max_f)$  is satisfiable then
3     | Add  $(\Phi \wedge cdt, l_t)$  to  $WL$ 
4   else if  $\Phi \wedge cdt \wedge (nb_f == max_f)$  is satisfiable then
5     | Stop injection in this path
6     | Add  $(\Phi \wedge cdt, l_t)$  to  $WL$ 
7   end
8   /* Idem for else branch ( $\neg cdt$ ) */
9 end

```

385 The first angle we explore to minimize query complexity is to reduce the
 386 number of injection points by *stopping the injection process as soon as possible*.
 387 Indeed, fewer injection points mean fewer extra symbolic variables and in general
 388 smaller and simpler queries for the SMT solver. We call this optimization *Early*
 389 *Detection of fault Saturation*, and write FASE-EDS when it is activated.

390 Its difference compared to FASE is in handling conditional jumps, illustrated
 391 in Algorithm 6. Instead of checking whether a branch can be explored without
 392 exceeding the maximum number of faults, we double the check: (1) first we check
 393 whether the branch can be explored with strictly fewer faults than allowed. If

394 the query is satisfiable, the analysis continues down that branch as usual; (2) if
 395 not satisfiable, we check whether the branch is feasible with exactly the maximal
 396 number of faults allowed. If not, the branch is infeasible and we stop as usual. Yet,
 397 if it is feasible, then we know that we have spent all allowed faults. We can thus
 398 continue the exploration *without injecting any new fault* in the corresponding
 399 search sub-tree, leading to simpler subsequent queries.

400 **Proposition 3.** *FASE-EDS is correct and k -complete for the adversarial reach-*
 401 *ability problem.*

402 *Proof.* FASE-EDS remains correct as it does not modify the path predicate
 403 computation, and it remains k -complete as it only prunes fault injections that
 404 are actually infeasible – and would have been proven so by the solver, later in
 405 the solving process.

406 5.5 Optimization via Injection on Demand (FASE-IOD)

407 The second angle explored to reduce query complexity through the reduction of
 408 injection points is *to inject faults on demand*, only when they are truly needed.
 409 We call this optimization *Injection On Demand*, and write FASE-IOD when it
 410 is activated.

411 To inject faults on demand, we now build *two* path predicate along a path:
 412 the working path predicate Φ based on which solver queries are built (where we
 413 try to minimize fault injection), and the normal adversarial path predicate Φ_F
 414 computed in previous sections (encompassing all the faults seen so far).

Algorithm 7: FASE-IOD assignment evaluation

Input: path predicate Φ , faulted path predicate Φ_F , assignment instruction

$x := expr$

Output: Updated Φ , Φ_F

```

1 Function eval_assign_IOD( $\Phi$ ,  $\Phi_F$ ,  $cdt$ ,  $x$ ,  $expr$ ) is
2   |  $\Phi_F$ ,  $expr' := \text{FaultEncoding}(\Phi_F, expr)$ 
3   | return ( $\Phi \wedge (x \triangleq expr)$ ,  $\Phi_F \wedge (x \triangleq expr')$ )
4 end
```

415 Algorithms are updated accordingly. Especially, assignment evaluation is du-
 416 plicated as shown in Algorithm 7: The normal symbolic assignment, with the
 417 original right-end-side expression $expr$, is added to Φ , while Φ_F is updated with
 418 the fault encoding of the assignment, $expr'$.

419 The on-demand reasoning takes place in the conditional jump instruction
 420 process detailed in Algorithm 8. The basic idea is to first check branch feasibility
 421 with the simpler path predicate Φ , encompassing the least number of faults. We
 422 continue this way as long as we can, meaning we rely on standard reachability
 423 as much as we can.

Algorithm 8: FASE-IOD conditional jump evaluation

Input: path predicate Φ , conditional jump instruction *if cdt then l_t else l_e*
Data: fault counter nb_f , maximal number of faults max_f , under approximation counter $under_counter$, worklist WL
Output: WL updated in place

```
1 Function eval_conditional_jump_IOD( $\Phi, \Phi_F, cdt, l_t, l_e$ ) is
2   if  $\Phi \wedge cdt \wedge (nb_f \leq max_f)$  is satisfiable then
3     | Add ( $\Phi \wedge cdt, \Phi_F \wedge cdt, l_t$ ) to  $WL$ 
4   else if  $under\_counter \leq max_f$  then
5     | if  $\Phi_F \wedge cdt \wedge (nb_f \leq max_f)$  is satisfiable then
6       |  $\Phi := \Phi_F$ 
7       |  $under\_counter := under\_counter + 1$ 
8       | Add ( $\Phi \wedge cdt, \Phi_F \wedge cdt, l_t$ ) to  $WL$ 
9     | end
10  end
11  /* Idem for else branch ( $\neg cdt$ ) */
12 end
```

424 When encountering a branch infeasible with Φ , we then check whether this
425 branch is feasible with all the possible faults seen so far, i.e. using Φ_F . If no
426 that is a stop, otherwise we know that Φ does not encompass enough faults to
427 go further. We then replace Φ by Φ_F (called a *switch*) at this stage, and thus
428 continue with strictly more faults. Note that this is straightforward as Φ_F and
429 Φ only differ on fault injections. Then again, the new Φ will not accumulate any
430 fault (until a new switch) while Φ_F continues accumulating all possible faults.

431 As a bonus, the number of path predicate switches gives us an under-
432 approximation $under_counter$ of the number of faults already needed in the
433 path under analysis. We use it to stop the injection early, when at least nb_f
434 faults have been used.

435 **Proposition 4.** *FASE-IOD is correct and k -complete for the adversarial reach-*
436 *ability problem.*

437 *Proof.* FASE-IOD explores the same feasible paths as FASE, hence preserving
438 its properties.

439 5.6 Optimizations combination

440 Both optimizations can be combined together as illustrated in Algorithm 9.
441 Taking FASE-IOD as a basis, saturation detection is added in the faulted path
442 predicate Φ_F queries at conditional branch handling. If the saturation is detected,
443 the main path predicate switch to Φ_F but Φ_F stops being updated and queried
444 further down that path, which stops fault injection.

445 **Proposition 5.** *The combination of FASE-EDS and FASE-IOD is correct and*
446 *k -complete for the adversarial reachability problem.*

447 *Proof.* This combination also explores all possible paths for the considered at-
448 tacker models, like FASE, hence preserving its properties.

Algorithm 9: FASE-IOD and FASE-EDS combination, conditional jump evaluation

Input: path predicate Φ , faulty path predicate Φ_F , conditional jump instruction *if cdt then l_t else l_e*

Data: fault counter nb_f , maximal number of faults max_f , under approximation counter $under_counter$, worklist WL

Output: WL updated in place

```

1 Function eval_conditional_jump_EDS_IOD( $\Phi, \Phi_F, cdt, l_t, l_e$ ) is
2   if  $\Phi \wedge cdt \wedge (nb_f \leq max_f)$  is satisfiable then
3     | Add  $(\Phi \wedge cdt, \Phi_F \wedge cdt, l_t)$  to  $WL$ 
4   else if  $under\_counter \leq max_f$  then
5     | if  $\Phi_F \wedge cdt \wedge (nb_f < max_f)$  is satisfiable then
6       |  $\Phi := \Phi_F$ 
7       |  $under\_counter := under\_counter + 1$ 
8       | Add  $(\Phi \wedge cdt, \Phi_F \wedge cdt, l_t)$  to  $WL$ 
9     | else if  $\Phi_F \wedge cdt \wedge (nb_f == max_f)$  is satisfiable then
10      |  $\Phi := \Phi_F$ 
11      | Stop  $\Phi'$  update and queries
12      | Add  $(\Phi \wedge cdt, \Phi_F \wedge cdt, l_t)$  to  $WL$ 
13    end
14  end
15 end
  
```

/* Idem for else branch ($\neg cdt$) */

449 6 Implementation

450 We now provide details about our forkless adversarial symbolic execution (FASE)
 451 implementation, named BINSEC/ASE, for Adversarial Symbolic Execution. *The*
 452 *code will be made open-source.*

453 **Binary-level fault injection.** While our method works for any program ab-
 454 straction level, we choose to implement it for the binary level, which makes more
 455 sense in many security scenarios. We implement our forkless adversarial symbolic
 456 execution on top of the BINSEC symbolic engine [10, 38, 40]. It has already been
 457 used in a number of significant case studies [9, 36, 37, 80, 81], and it is notably
 458 able to achieve bounded verification (k-completeness) and to reasonably deal
 459 with symbolic pointers [44].

460 We modified the path predicate computation of BINSEC 0.4.0 as described
 461 in section 5, and implemented our dedicated optimizations FASE-EDS, FASE-
 462 IOD and FASE EDS+IOD. The attacker goal is specified as a local predicate
 463 to reach, using BINSEC directives. We currently support data faults such as
 464 arbitrary modification, bit-flip, and reset of variables. Test inversion is emulated
 465 through faulting the condition of conditional jumps. We let the user define an in-
 466 jection target range, made of multiple code address intervals. For large programs,
 467 it enables focusing on the security critical sections. Finally, we also provide a

468 blacklist for some memory locations which will never be faulted. The blacklist is
 469 mostly used for the stack register (`esp` in x86, which is concretized in the anal-
 470 ysis) and the program counter, as our fault model does not include tampering
 471 with the stack nor arbitrary control faults.

472 **Details.** Our exploration strategy is depth first, the underlying SMT solver is
 473 Bitwuzla [71]. We constrain the faulted values to differ from the original values
 474 in fault encodings, such that only true corruptions are reported as active faults.

475 7 Evaluation

476 We now evaluate our new algorithm for software verification against multi-fault
 477 attacks. We consider the following research questions.

- 478 – **RQ1:** is our tool correct and complete? In particular, can we find attacks
- 479 on vulnerable programs and prove secure resistant programs?
- 480 – **RQ2:** can we scale in number of faults without path explosion?
- 481 – **RQ3:** what is the impact of our optimizations?

482 Besides this evaluation, we also show the use of our method in a number of
 483 different security scenarios (Section 7.5), and on a larger case study (Section 8).

484 7.1 Experimental setting

485 **The machine used.** We ran our experiments on a cloud machine with a proces-
 486 sor Intel Dual Xeon 4214R with 48 CPU cores and 384GB of RAM. Experiments
 487 ran in parallel on the 48 cores, each run using only one core.

488 **The attacker model** chosen in this evaluation can perform a varying number
 489 of faults. Its goal is expressed as a security oracle directly written in C for each
 490 benchmark, the computation of which is not faulted.

491 **The benchmark** used here is a standard set of programs from the SWiFI
 492 literature on physical fault injections and high-security devices, characterized
 493 in Table 2. First, the 8 versions of VerifyPIN from the FISSC [42] benchmark
 494 suite, dedicated to the evaluation of physical fault attack analyses. VerifyPIN is
 495 an authentication program. There are one unprotected and 7 different protected
 496 versions, some vulnerable, some resistant to one test inversion fault. We added
 497 2 manually unrolled versions of the unprotected VerifyPIN, with a PIN size of
 498 4 and 16, to add diversity in the benchmarks with programs without loops. An
 499 oracle is provided by FISSC, checking if the user PIN truly corresponds to the
 500 reference PIN. Second, we take the 2 versions of the npo2 program from Le
 501 et al. work [65], together with their oracles. Npo2 is a program computing an
 502 integer’s upper power of two. The attacker’s goal is to perform a silent data
 503 corruption, i.e. change the end result without triggering countermeasures. One
 504 version is vulnerable to one arbitrary data fault, the second is resistant due to
 505 extra arithmetic checks.

506 **Compilation.** The benchmarks are written in C and have been compiled with
507 gcc for the Intel x86-32 architecture, using the flag “-O0” to preserve counter-
508 measures. For BINSEC compatibility, we use the “-static” flag to include the
509 necessary library functions directly in the binary.

Table 2: Benchmarks characteristics and statistics of a standard SE analysis

Program group (#)	BINSEC analysis - no fault					
	C loc	x86 loc	#instruction (explored)	#paths	#branch in a path	Time
Section 7						
VerifyPINs (8)	80-140	160-215	192-269	1	17-34	< 0.1s
VerifyPIN unrolled (2)	40-85	140-430	142-442	5-17	5-17	< 0.1s
np02 (2)	50	200-220	607-653	3	31-33	< 0.1s
Section 8						
WooKey bootloader	3.2k	2350	290k	17	18k	9s
Section 7.5						
CRT-RSA (3)	125-170	400-600	108k-29M	1	5k-1.3M	0.4s - 1m27
Secret keeping machine (2)	100-200	240-360	1k-1.3k	1	130-150	< 0.1s
VerifyPIN_0 with SecSwift	80	430	430	1	22	< 0.1s

510 **BINSEC settings.** We limit the maximal depth of an analysis to the depth
511 necessary to perform an exhaustive non-faulty analysis, rounded to the upper
512 hundred. We exhaustively explore all the possible paths up to this bound and do
513 not stop at the first identified attack, in order to have comparable results. We
514 set the global analysis timeout for 1 day. We fault values and not addresses, we
515 do not directly fault the stack pointer or the program counter, and we do not
516 fault the status flags unless explicitly specified.

517 7.2 Correctness and completeness in practice (RQ1)

518 We first show that our tool works as expected on several codes with known
519 ground truth. (1) We check that indeed, with no fault allowed, no attack is
520 found in any of the benchmarks; (2) We check that indeed the insecure np02
521 program is vulnerable to a single arbitrary data fault while the secure version is
522 not – it can still be exploited with two faults; (3) According to their authors, the
523 VerifyPIN versions 0 to 4 are vulnerable to one test inversion, while VerifyPIN
524 5 to 7 are resistant to it. We indeed reproduce these results. When allowing two
525 faults, all VerifyPIN become vulnerable; (4) When using one arbitrary data fault
526 against the VerifyPINs, all versions are found vulnerable. We manually check
527 that indeed the identified attack paths make sense; (5) Our manually unrolled
528 versions of VerifyPINs do not contain conditional branching instructions in the

529 targeted function, making them resistant to test inversion. We check that this is
530 the case, while they are still vulnerable to a single arbitrary data fault.

531 **Conclusion.** Our tool indeed can showcase a program vulnerability to fault
532 injection attacks and prove resistance to fault injection attacks, as expected by
533 the correctness and k-completeness properties of the underlying algorithms.

534 7.3 Scalability (RQ2)

535 For this evaluation, we focus on an attacker capable of arbitrary data faults, as
536 those weigh the heaviest on the analysis.

537 We take FASE-IOD as our best performing technique (see Section 7.4). We
538 evaluate here its capability to handle multi-fault and avoid path explosion, com-
539 pared to the forking technique. Results are illustrated in Figures 4 and 5. Note
540 that all FASE variants explore the same number of paths, and are thus repre-
541 sented as FASE in Figure 5. For each benchmark, we took the arithmetic mean
542 for 100 runs. Values presented here are the geometric mean over the benchmarks.
543 More results are available in Appendix B, Table 4.

544 FASE-IOD is 10x times faster than Forking for 1 fault, and x200 times faster
545 for 2 faults on average. For the best case benchmark, we are x224 times faster for
546 1 fault and x6121 for 2. Starting from three faults onward, Forking experiences
547 timeouts, rendering values non comparable. Half of the benchmark timeouts for
548 3 faults, three quarters for 4 faults, 11 over 12 for 6 faults and all of them after
549 that. FASE-IOD never timeouts in this experiment. This scaling is enabled by
550 avoiding path explosion. On average, Forking explores x50 times more paths for
551 2 faults than for one, while FASE-IOD only explores x3 times more paths. From
552 Figure 4, we see FASE on its own already scales better than Forking, being
553 x3 times faster for 1 fault and x108 times faster for 2, and never experiencing
554 timeouts either.

555 **Conclusion.** FASE-IOD shows improved scalability in terms of the maximum
556 number of faults allowed, for the arbitrary data fault model, compared to the
557 forking technique.

558 7.4 Performance optimization (RQ3)

559 We evaluate our forkless variants: FASE, FASE-EDS, FASE-IOD and FASE
560 EDS+IOD, to determine which performs best for arbitrary data faults. Results
561 are illustrated in Figures 4, 5 and 6. Raw results are available in Appendix B,
562 Tables 4 and 5.

563 We vary again the maximum number of faults from 1 to 10. Note that all
564 FASE variants explore the same number of paths for each number of faults, as
565 the optimizations reduce the number of faults injected but do not lose correctness
566 nor k-completeness. FASE indeed generates complex queries, taking on average
567 around twice the time necessary for Forking queries to be solved. FASE-EDS
568 then gains a little bit in that regard. FASE’s queries take only x1.04 longer to




Fig. 4: Analysis time

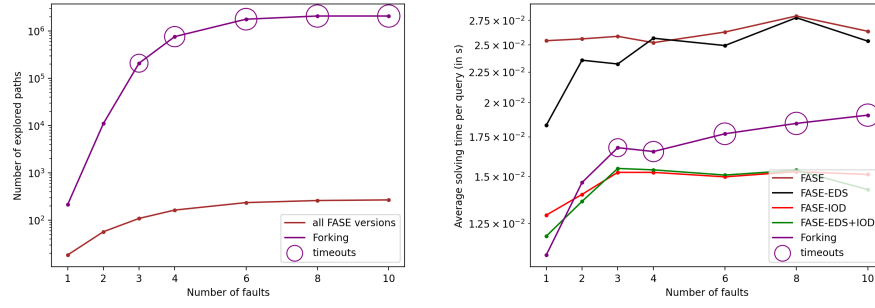


Fig. 5: Average number of explored paths, Average solving time per query




Fig. 6: Number of queries sent to the solver

569 solve on average for all fault numbers. The real improvement comes with the On-
570 Demand logic of FASE-IOD (x2.02 times faster on average over all fault numbers)
571 and FASE EDS+IOD (x2.02 also), where query complexity drops to the level
572 of Forking. This improvement in query complexity is achieved algorithmically
573 through more queries created. However, due to more queries being arithmetically
574 simplified, fewer queries are sent to the solver for FASE-IOD (x0.88 on average
575 over all fault values compared with FASE) and FASE EDS+IOD (x0.98). FASE-
576 EDS sent approximately the same number of queries as FASE. The number of
577 queries sent to the solver explodes for Forking, correlated with the path explosion
578 experienced. In terms of performance, two trends appear as the number of faults
579 allowed increases. FASE and FASE-EDS tend to be between x2 and x3 times
580 slower than FASE-IOD and FASE EDS+IOD. In the end, FASE-IOD proves to
581 be the fastest optimization (x1.1 times faster than FASE EDS+IOD on average
582 over all number of faults), likely due to the combination of on-demand logic and
583 fewer queries than FASE EDS+IOD.

584 **Conclusion.** We retain FASE-IOD as our best performing forkless adversarial
585 algorithm, at most x3.06 faster than FASE.

586 7.5 Other experiments and fault models

587 **CRT-RSA.** Puys et al. [78] describe three versions of CRT-RSA: unprotected,
588 Shamir version and Aumuller version. Only the last one is shown to resist the
589 BellCoRe attack [16] which uses a single reset fault to break the cryptography.
590 We were able to automatically reproduce the attack with 1 reset fault on the
591 unprotected version of CRT-RSA, after 3s of analysis, and we were not able to
592 find attacks on the other two versions in 10 days time.

593 **Secret-keeping machine.** Dullien [41] proposes two versions of a secret-keeping
594 machine. The one based on linked lists is manually shown to be exploitable by
595 an attacker able to perform a single bit-flip, while the array version is shown to
596 be secure against that. For this benchmark, we activated faults on variables used
597 as addresses. We were able to reproduce the attack on our linked list implemen-
598 tation with one bit-flip fault. However, our implementation of the array version
599 proved to be vulnerable to a bit-flip in array indices, overwriting the victim's
600 password with the attacker's, enabling them to retrieve the victim's secret.

601 **SecSwift countermeasure.** We applied the SecSwift countermeasure, a binary-
602 level protection developed by STMicroelectronics [27, 45], to VerifyPIN version
603 0. We were able to find attacks on this binary with either a single test inversion
604 or a single arbitrary data fault.

605 8 Case study: the WooKey bootloader

606 We now confront our tool to a real-life security system, WooKey.

607 **Presentation of WooKey.** First presented in 2018 by ANSSI, the French sys-
608 tem security agency, the WooKey platform [14, 88] is “a custom STM32-based

609 USB thumb drive with mass storage capabilities designed for user data encryption and protection, with a full-fledged set of in-depth security defenses”. Their
610 choice to be open source and open hardware makes WooKey a relevant case
611 study: it is a real-life, complex device, security focused and available for reproducibility. Note also that Wookey has been extensively analyzed, as it was the
612 target of an ANSSI cybersecurity challenge for security professionals [5].
613
614

615 **Security scenario and goal of our study.** We focus on WooKey bootloader, a dual-bank system enabling hot firmware updates. The system is hardened, especially redundant test protections are present in critical sections to protect
616 against test inversion faults. We consider the same attacker model as the ANSSI challenge did [5]: the attacker seeks to manipulate the bootloader logic to boot
617 on the older firmware, more likely to contain security vulnerabilities. We also consider an attacker able to perform a single arbitrary data fault. We see in
618 Table 2 that WooKey bootloader size is orders of magnitude larger than the programs used for evaluation in Section 7. Wookey is available as C code. We
619 compile it like we did for the evaluation benchmarks (Section 7.1).
620
621
622
623
624

625 We conduct the following three analyses:

- 626 1. automatically analyze WooKey at binary-level to check whether we are able
627 to find previously known faults [63], and/or new ones: we are indeed able
628 to find the two faults identified by prior work [63] (A1, A2), *as well as a*
629 *completely new one* (A3);
- 630 2. automatically analyze at binary-level the patch version of Wookey proposed
631 by Lacombe et al. [63]: we found that the proposed patch indeed blocks the
632 two known attacks (A1 and A2), but not the new attack (A3);
- 633 3. propose a definitive patch by adding a counter-measure for A3 and remove
634 parts of the counter-measures which are shown to be useless here. The patch
635 is proven correct w.r.t. our attack model.

636 We discuss these results in the following (see also Appendix C) and we present
637 briefly in Section 8 the discovery of two more known faults. Overall, it demon-
638 strates that our technique can scale to binary-level real-size systems.

639 **Analyze key parts of Wookey.** Lacombe et al. find an attack in the *loader_*
640 *exec_req_selectbank* function (A1) and another in the *loader_exec_req_*
641 *flashlock* function (A2). They correspond to data corruption in branching condi-
642 tions. We are able to find both attacks, linking faults back to their locations in the
643 C code with debug information. *We also find an additional attack* not mentioned
644 in Lacombe et al., faulting another part of the *loader_exec_req_flashlock*
645 function (A3).

646 **Analyze a security patch of WooKey.** We now evaluate the protection
647 scheme proposed by Lacombe et al. [63] for these attacks. It consists of four
648 extra counter-measures named from C1 to C4. We found indeed that the full
649 protection prevents attacks A1 and A2, as claimed by the authors of the patch.
650 Yet, our analysis shows that the protection does not prevent the new attack A3.

651 **Propose a new patch and evaluate it.** We manually inspect these different
652 analysis results to understand what happens. We have especially been able to

653 identify the root cause of A3 and propose a dedicated countermeasure for it
654 (named CMA). Also, by analyzing each counter-measure in isolation, we have
655 been able to understand that counter-measures CM1 and CM3 do not block any
656 attack path as they are redundant with other tests in the code and can be safely
657 removed. Overall, our new patch (CMA + refined former patch) is shown by
658 our tool to protect against all the attacks, for an attacker able to perform one
659 arbitrary data fault.

Table 3: Table summarizing the effects of countermeasures

Protection scheme	A1	A2	A3 (new)
	1.3	1.31	1.25
Normal Wookey	✓	✓	✓
Prior patch (CM1+CM2+CM3+CM4)	✗	✗	✓
Our patch (CM2+CM4+CMA)	✗	✗	✗

Legend - ✓: attack path found by our tool / ✗: no attack found

660 **Other attacks on WooKey.** We were also able to find two other known at-
661 tacks on Wookey. (**Attack vector combination**) The iso8716 library, used in
662 WooKey for secure communication, presents a vulnerability to fault injection
663 which enables a software buffer-overflow in function *SC_get_ATR* [63]. Us-
664 ing an attacker with a single arbitrary data fault, we were able to reproduce
665 this attack; (**Faulty redundant test**) Martin et al. [68] shows an incorrect im-
666 plementation of a redundant test to prevent single test inversion faults in the
667 *loader_set_state* function. We reproduce this result.

668 9 Discussion

669 **Fault models.** Our current approach does not support advanced *control* faults
670 such as instruction corruption or instruction skip. Instruction corruption is out
671 of scope as it permanently changes an instruction, while we modify computa-
672 tion results. It is related to self-modification, a notoriously difficult point to
673 address in adversarial binary-level code analysis [17, 77]. Instruction skip (or
674 other arbitrary control jumps) could be modeled by local modification of the
675 program counter, yet at the price of a huge path explosion. Also, regarding
676 micro-architectural attacks, while modeling Meltdown-style effects is straight-
677 forward (bit-flips), modeling Spectre attacks is much more difficult due to the
678 speculative windows mechanism.

679 **Other formal methods.** While in the paper we focus on symbolic execution, we
680 believe the main optimization ideas developed here can be used with other formal
681 techniques, e.g. Bounded Model Checking [29, 31], Abstract Interpretation [34]

682 or CEGAR model checking [30]. Note that for each of them, fault injection may
683 result either in path explosion or precision loss. Still, our forkless encoding should
684 be able to help at least all approaches based to some extent on path unrolling.

685 **Other properties.** The forkless encoding can surely benefit other classes of
686 properties to be achieved by the attacker, especially those known to be sup-
687 ported by (extensions of) symbolic execution, for example: trace properties such
688 as use-after-free, k-hyperreachability properties (secret leakage, privacy leakage,
689 violation of constant-time, etc.) [36], the recent robust reachability proposal [48]
690 for replicable bugs, etc. Our formalism itself is quite generic and can accom-
691 modate a wide range of properties, as we mainly keep the property unchanged
692 but modify the underlying transition system. We could for example imagine an
693 attacker willing to activate a non-terminating execution (denial of service).

694 **Forkless encoding and instrumentation.** Several prior works use code-level
695 instrumentation [68] or llvm-level instrumentation [63, 65, 76] in order to leverage
696 standard program analyzers as is. The forkless encoding we propose can also be
697 used this way, for more flexibility but without additional optimizations. Actually,
698 we performed some experiments with Klee and a C-level forkless instrumenta-
699 tion, and do observe significant improvement.

700 10 Related Work

701 **SWiFI.** Prior work in SWiFI has already been discussed in Section 3. All meth-
702 ods in this domain consider low-level formalism: C [28, 68], llvm [63, 76], bi-
703 nary [15, 20, 25, 50]. Half of the techniques rely on the mutant approach [25, 28, 49,
704 50, 79], and the other half relies on forking [15, 20, 63, 76]. While most approaches
705 target attack finding (with symbolic execution and bounded model-checking),
706 some do aim at full verification [79], especially with deductive verification [28, 68].
707 Very few works consider multi-faults [63, 68, 76]. Interestingly, Lacombe et al. [63]
708 propose a static way of reducing injection points on C programs, that is com-
709 plementary to our own method – still, static analysis at binary-level is known to
710 be hard. Note that a few methods do consider instruction skips [20, 49, 50], yet
711 with path explosion issues.

712 **Robustness analysis.** SWiFI is also used for robustness evaluation [32, 56,
713 64, 65, 72, 74, 87, 89], in order to verify the correct behavior of error handling
714 mechanisms. They rely also on forking techniques. The fault models are similar
715 to hardware fault injection, yet multi-fault is not really an issue there, as faults
716 are supposed to originate from safety issues (e.g. cosmic rays) and have no reason
717 to accumulate unreasonably.

718 **Formalizations and fault models.** While it is common in the field of au-
719 tomated formal verification of cryptographic protocols to consider models of
720 attackers (typically, extensions of the “Dolev-Yao” model) – either by specifying
721 what the attackers can do [2] or what they cannot do [7], only very few for-
722 malizations of software-level attacker capabilities have been proposed so far. In

software security, control-flow integrity attacks have been categorized by the capability an attacker needs [21], but these efforts have been restricted to manual reasoning. Interestingly, Given-Wilson et al. [51] propose a formalization of fault injection using Turing machines, but to our knowledge, no algorithm has been built for it. Also, Fournet et al. [46] propose a type system for program-level non-interference, taking into account an active adversary modeled as adversarial components able to perform any action at certain steps of the program.

Mutation testing. Sometimes called software fault injection, mutation testing [33, 75] aims to generate a comprehensive test suite by building test cases discriminating various mutants of a program, and is recognized as a very powerful testing criterion. As it focuses on coverage, mutant explosion cannot be avoided. Dedicated SE techniques [8, 11, 67, 73] have been designed.

11 Conclusion

We formalize the concept of adversarial reachability, extending standard reachability to include the presence of an advanced attacker in program analysis, and we propose a dedicated symbolic algorithm for adversarial reachability, integrating a novel forkless encoding of faults together with dedicated optimizations. Our technique is shown to significantly reduce the number of paths to explore, and scales up to 10 faults on a standard SWiFI benchmark, where prior forking attempts timeout for 3 faults. Also, we show that our method scale to realistic size examples, such as the WooKey project where we have been able to replay known fault attacks and to even find a novel vulnerability in a recently proposed countermeasure patch [63].

References

1. Abadi, M., Budiu, M., Erlingsson, U., Ligatti, J.: Control-flow integrity principles, implementations, and applications. *ACM Transactions on Information and System Security (TISSEC)* **13**(1), 1–40 (2009)
2. Abadi, M., Cortier, V.: Deciding knowledge in security protocols under equational theories. In: *International Colloquium on Automata, Languages, and Programming*. pp. 46–58. Springer (2004)
3. Akhunzada, A., Sookhak, M., Anuar, N.B., Gani, A., Ahmed, E., Shiraz, M., Furnell, S., Hayat, A., Khan, M.K.: Man-at-the-end attacks: Analysis, taxonomy, human aspects, motivation and future directions. *Journal of Network and Computer Applications* **48**, 44–57 (2015)
4. Anceau, S., Bleuet, P., Clédière, J., Maingault, L., Rainard, J.I., Tucoulou, R.: Nanofocused X-ray beam to reprogram secure circuits. In: *International Conference on Cryptographic Hardware and Embedded Systems*. pp. 175–188. Springer (2017)
5. ANSSI, Amossys, EDSI, LETI, Lexfo, Oppida, Quarkslab, SERMA, Synacktiv, Thales, Labs, T.: Inter-cesti: Methodological and technical feedbacks on hardware devices evaluations. In: *SSTIC 2020, Symposium sur la sécurité des technologies de l’information et des communications* (2020)

- 764 6. Ball, T., Cook, B., Levin, V., Rajamani, S.K.: Slam and static driver verifier: Tech-
765 nology transfer of formal methods inside microsoft. In: International Conference
766 on Integrated Formal Methods. pp. 1–20. Springer (2004)
- 767 7. Bana, G., Comon-Lundh, H.: A computationally complete symbolic attacker for
768 equivalence properties. In: Proceedings of the 2014 ACM SIGSAC Conference on
769 Computer and Communications Security. pp. 609–620 (2014)
- 770 8. Bardin, S., Chebaro, O., Delahaye, M., Kosmatov, N.: An all-in-one toolkit for
771 automated white-box testing. In: International Conference on Tests and Proofs.
772 pp. 53–60. Springer (2014)
- 773 9. Bardin, S., David, R., Marion, J.Y.: Backward-bounded dse: targeting infeasibility
774 questions on obfuscated codes. In: 2017 IEEE Symposium on Security and Privacy
775 (SP). pp. 633–651. IEEE (2017)
- 776 10. Bardin, S., Herrmann, P., Leroux, J., Ly, O., Tabary, R., Vincent, A.: The bincoa
777 framework for binary code analysis. In: International Conference on Computer
778 Aided Verification. pp. 165–170. Springer (2011)
- 779 11. Bardin, S., Kosmatov, N., Cheynier, F.: Efficient leveraging of symbolic execution
780 to advanced coverage criteria. In: 2014 IEEE Seventh International Conference on
781 Software Testing, Verification and Validation. pp. 173–182. IEEE (2014)
- 782 12. Barrett, C., Tinelli, C.: Satisfiability modulo theories. In: Handbook of model
783 checking, pp. 305–343. Springer (2018)
- 784 13. Barthe, G., Dupressoir, F., Fouque, P.A., Grégoire, B., Zapalowicz, J.C.: Synthesis
785 of fault attacks on cryptographic implementations. In: Proceedings of the 2014
786 ACM SIGSAC Conference on Computer and Communications Security. pp. 1016–
787 1027 (2014)
- 788 14. Benadjila, R., Renard, M., Trebuchet, P., Thierry, P., Michelizza, A., Lefaure, J.:
789 Wookey: Usb devices strike back. Proceedings of SSTIC (2018)
- 790 15. Berthier, M., Bringer, J., Chabanne, H., Le, T.H., Rivière, L., Servant, V.: Idea:
791 embedded fault injection simulator on smartcard. In: International Symposium on
792 Engineering Secure Software and Systems. pp. 222–229. Springer (2014)
- 793 16. Boneh, D., DeMillo, R.A., Lipton, R.J.: On the importance of checking crypto-
794 graphic protocols for faults. In: International conference on the theory and appli-
795 cations of cryptographic techniques. pp. 37–51. Springer (1997)
- 796 17. Bonfante, G., Fernandez, J., Marion, J.Y., Rouxel, B., Sabatier, F., Thierry, A.:
797 Codisasm: Medium scale concat disassembly of self-modifying binaries with over-
798 lapping instructions. In: Proceedings of the 22nd ACM SIGSAC Conference on
799 Computer and Communications Security. pp. 745–756 (2015)
- 800 18. Bounimova, E., Godefroid, P., Molnar, D.: Billions and billions of constraints:
801 Whitebox fuzz testing in production. In: 2013 35th International Conference on
802 Software Engineering (ICSE). pp. 122–131. IEEE (2013)
- 803 19. Bozzato, C., Focardi, R., Palmarini, F.: Shaping the glitch: optimizing voltage fault
804 injection attacks. IACR Transactions on Cryptographic Hardware and Embedded
805 Systems pp. 199–224 (2019)
- 806 20. Bréjon, J.B., Heydemann, K., Encrenaz, E., Meunier, Q., Vu, S.T.: Fault attack
807 vulnerability assessment of binary code. In: Proceedings of the Sixth Workshop on
808 Cryptography and Security in Computing Systems. pp. 13–18 (2019)
- 809 21. Burow, N., Carr, S.A., Nash, J., Larsen, P., Franz, M., Brunthaler, S., Payer,
810 M.: Control-flow integrity: Precision, security, and performance. ACM Computing
811 Surveys (CSUR) **50**(1), 1–33 (2017)
- 812 22. Cadar, C., Dunbar, D., Engler, D.R., et al.: Klee: unassisted and automatic gen-
813 eration of high-coverage tests for complex systems programs. In: OSDI. vol. 8, pp.
814 209–224 (2008)

- 815 23. Cadar, C., Ganesh, V., Pawlowski, P.M., Dill, D.L., Engler, D.R.: Exe: Automat-
816 ically generating inputs of death. *ACM Transactions on Information and System*
817 *Security (TISSEC)* **12**(2), 1–38 (2008)
- 818 24. Cadar, C., Sen, K.: Symbolic execution for software testing: three decades later.
819 *Communications of the ACM* **56**(2), 82–90 (2013)
- 820 25. Carré, S., Desjardins, M., Facon, A., Guilley, S.: Openssl bellcore's protection helps
821 fault attack. In: 2018 21st Euromicro Conference on Digital System Design (DSD).
822 pp. 500–507. IEEE (2018)
- 823 26. Cervesato, I.: The dolev-yao intruder is the most powerful attacker. In: 16th Annual
824 Symposium on Logic in Computer Science—LICS. vol. 1, pp. 1–2. Citeseer (2001)
- 825 27. Chauvet, H., de Ferrière, F., Bizet, T.: Software fault injection for secswift quali-
826 fication (2021), <https://jaif.io/2021/media/JAIF2021%20-%20deFerriere.pdf>
- 827 28. Christofi, M., Chetali, B., Goubin, L.: Formal verification of an implementation of
828 crt-rsa vigilant's algorithm. In: PROOFS workshop: pre-proceedings. vol. 28 (2013)
- 829 29. Clarke, E., Biere, A., Raimi, R., Zhu, Y.: Bounded model checking using satisfia-
830 bility solving. *Form. Methods Syst. Des.* (2001)
- 831 30. Clarke, E., Grumberg, O., Jha, S., Lu, Y., Veith, H.: Counterexample-guided ab-
832 straction refinement for symbolic model checking. *J. ACM* (2003)
- 833 31. Clarke, E., Kroening, D., Lerda, F.: A tool for checking ansi-c programs. In: Inter-
834 national Conference on Tools and Algorithms for the Construction and Analysis
835 of Systems. pp. 168–176. Springer (2004)
- 836 32. Cotroneo, D., De Simone, L., Liguori, P., Natella, R.: Profipy: Programmable soft-
837 ware fault injection as-a-service. In: 2020 50th annual IEEE/IFIP international
838 conference on dependable systems and networks (DSN). pp. 364–372. IEEE (2020)
- 839 33. Cotroneo, D., Natella, R.: Fault injection for software certification. *IEEE Security*
840 *& Privacy* **11**(4), 38–45 (2013)
- 841 34. Cousot, P.: Abstract interpretation. *ACM Computing Surveys (CSUR)* **28**(2), 324–
842 328 (1996)
- 843 35. Cousot, P., Cousot, R., Feret, J., Mauborgne, L., Miné, A., Monniaux, D., Rival,
844 X.: The astreé analyzer. In: Programming Languages and Systems (2005)
- 845 36. Daniel, L.A., Bardin, S., Rezk, T.: Binsec/rel: Efficient relational symbolic execu-
846 tion for constant-time at binary-level. In: 2020 IEEE Symposium on Security and
847 Privacy (SP). pp. 1021–1038. IEEE (2020)
- 848 37. Daniel, L.A., Bardin, S., Rezk, T.: Hunting the haunter-efficient relational symbolic
849 execution for spectre with haunted relse. In: NDSS (2021)
- 850 38. David, R., Bardin, S., Ta, T.D., Mounier, L., Feist, J., Potet, M.L., Marion,
851 J.Y.: Binsec/se: A dynamic symbolic execution toolkit for binary-level analysis.
852 In: SANER (2016)
- 853 39. Dehbaoui, A., Dutertre, J.M., Robisson, B., Tria, A.: Electromagnetic transient
854 faults injection on a hardware and a software implementations of AES. In: 2012
855 Workshop on Fault Diagnosis and Tolerance in Cryptography. pp. 7–15. IEEE
856 (2012)
- 857 40. Djoudi, A., Bardin, S.: Binsec: Binary code analysis with low-level regions. In: In-
858 ternational Conference on Tools and Algorithms for the Construction and Analysis
859 of Systems. pp. 212–217. Springer (2015)
- 860 41. Dullien, T.: Weird machines, exploitability, and provable unexploitability. *IEEE*
861 *Transactions on Emerging Topics in Computing* **8**(2), 391–403 (2017)
- 862 42. Dureuil, L., Petiot, G., Potet, M.L., Le, T.H., Crohen, A., Choudens, P.d.: Fissc:
863 A fault injection and simulation secure collection. In: International Conference on
864 Computer Safety, Reliability, and Security. pp. 3–11. Springer (2016)

- 865 43. Facebook: Infer static analyzer. <https://fbinfer.com/>
- 866 44. Farinier, B., David, R., Bardin, S., Lemerre, M.: Arrays made simpler: An efficient,
867 scalable and thorough preprocessing. In: LPAR. pp. 363–380 (2018)
- 868 45. de Ferrière, F.: Software countermeasures in the llvm risc-v compiler (2021),
869 [https://open-src-soc.org/2021-03/media/slides/3rd-RISC-V-Meeting-2021-03-30-](https://open-src-soc.org/2021-03/media/slides/3rd-RISC-V-Meeting-2021-03-30-15h00-Fran%C3%A7ois-de-Ferri%C3%A8re.pdf)
870 [15h00-Fran%C3%A7ois-de-Ferri%C3%A8re.pdf](https://open-src-soc.org/2021-03/media/slides/3rd-RISC-V-Meeting-2021-03-30-15h00-Fran%C3%A7ois-de-Ferri%C3%A8re.pdf)
- 871 46. Fournet, C., Rezk, T.: Cryptographically sound implementations for typed
872 information-flow security. In: Necula, G.C., Wadler, P. (eds.) Proceedings of the
873 35th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Lan-
874 guages, POPL 2008, San Francisco, California, USA, January 7-12, 2008. ACM
875 (2008)
- 876 47. Gangolli, A., Mahmoud, Q.H., Azim, A.: A systematic review of fault injection
877 attacks on iot systems. *Electronics* **11**(13), 2023 (2022)
- 878 48. Girol, G., Farinier, B., Bardin, S.: Not all bugs are created equal, but robust
879 reachability can tell the difference. In: International Conference on Computer Aided
880 Verification. pp. 669–693. Springer (2021)
- 881 49. Given-Wilson, T., Jafri, N., Lanet, J.L., Legay, A.: An automated formal process
882 for detecting fault injection vulnerabilities in binaries and case study on present.
883 In: 2017 IEEE Trustcom/BigDataSE/ICSS. pp. 293–300. IEEE (2017)
- 884 50. Given-Wilson, T., Jafri, N., Legay, A.: Combined software and hardware fault
885 injection vulnerability detection. *Innovations in Systems and Software Engineering*
886 **16**(2), 101–120 (2020)
- 887 51. Given-Wilson, T., Legay, A.: Formalising fault injection and countermeasures. In:
888 Proceedings of the 15th International Conference on Availability, Reliability and
889 Security. pp. 1–11 (2020)
- 890 52. Godefroid, P., Klarlund, N., Sen, K.: Dart: Directed automated random testing.
891 In: Proceedings of the 2005 ACM SIGPLAN conference on Programming language
892 design and implementation. pp. 213–223 (2005)
- 893 53. Godefroid, P., Levin, M.Y., Molnar, D.: Sage: whitebox fuzzing for security testing.
894 *Communications of the ACM* **55**(3), 40–44 (2012)
- 895 54. Goyal, B., Sitaraman, S., Venkatesan, S.: A unified approach to detect binding
896 based race condition attacks. In: Int’l Workshop on Cryptology & Network Security
897 (CANS). p. 16 (2003)
- 898 55. Gravelier, J., Dutertre, J.M., Teglia, Y., Moundi, P.L.: Faultline: Software-based
899 fault injection on memory transfers. In: 2021 IEEE International Symposium on
900 Hardware Oriented Security and Trust (HOST). pp. 46–55. IEEE (2021)
- 901 56. Hari, S.K.S., Tsai, T., Stephenson, M., Keckler, S.W., Emer, J.: Sassifi: An
902 architecture-level fault injection tool for gpu application resilience evaluation. In:
903 2017 IEEE International Symposium on Performance Analysis of Systems and
904 Software (ISPASS). pp. 249–258. IEEE (2017)
- 905 57. Van den Herrewegen, J., Oswald, D., Garcia, F.D., Temeiza, Q.: Fill your boots:
906 Enhanced embedded bootloader exploits via fault injection and binary analysis.
907 *IACR Transactions on Cryptographic Hardware and Embedded Systems* pp. 56–
908 81 (2021)
- 909 58. Karaklajić, D., Schmidt, J.M., Verbauwheide, I.: Hardware designer’s guide to
910 fault attacks. *IEEE Transactions on Very Large Scale Integration (VLSI) Systems*
911 **21**(12), 2295–2306 (2013)
- 912 59. Kim, C.H., Quisquater, J.J.: Fault attacks for CRT based RSA: New attacks, new
913 results, and new countermeasures. In: IFIP International Workshop on Information
914 Security Theory and Practices. pp. 215–228. Springer (2007)

- 915 60. Kirchner, F., Kosmatov, N., Prevosto, V., Signoles, J., Yakobowski, B.: Frama-c:
916 A software analysis perspective. *Form. Asp. Comput.* (2015)
- 917 61. Kocher, P., Horn, J., Fogh, A., Genkin, D., Gruss, D., Haas, W., Hamburg, M.,
918 Lipp, M., Mangard, S., Prescher, T., et al.: Spectre attacks: Exploiting speculative
919 execution. In: *SP* (2019)
- 920 62. Kocher, P., Horn, J., Fogh, A., Genkin, D., Gruss, D., Haas, W., Hamburg, M.,
921 Lipp, M., Mangard, S., Prescher, T., et al.: Spectre attacks: Exploiting speculative
922 execution. *Communications of the ACM* **63**(7), 93–101 (2020)
- 923 63. Lacombe, G., Feliot, D., Boespflug, E., Potet, M.L.: Combining static analysis
924 and dynamic symbolic execution in a toolchain to detect fault injection vulner-
925 abilities. In: *PROOFS WORKSHOP (SECURITY PROOFS FOR EMBEDDED*
926 *SYSTEMS)* (2021)
- 927 64. Larsson, D., Hähnle, R.: Symbolic fault injection. In: *International Verification*
928 *Workshop (VERIFY)*. vol. 259, pp. 85–103. Citeseer (2007)
- 929 65. Le, H.M., Herdt, V., Große, D., Drechsler, R.: Resilience evaluation via symbolic
930 fault injection on intermediate code. In: *2018 Design, Automation & Test in Europe*
931 *Conference & Exhibition (DATE)*. pp. 845–850. IEEE (2018)
- 932 66. Le, Q.L., Raad, A., Villard, J., Berdine, J., Dreyer, D., O’Hearn, P.W.: Finding
933 real bugs in big programs with incorrectness logic. *Proceedings of the ACM on*
934 *Programming Languages* **6**(OOPSLA1), 1–27 (2022)
- 935 67. Marcozzi, M., Bardin, S., Kosmatov, N., Papadakis, M., Prevosto, V., Correnson,
936 L.: Time to clean your test objectives. In: *Proceedings of the 40th International*
937 *Conference on Software Engineering*. pp. 456–467 (2018)
- 938 68. Martin, T., Kosmatov, N., Prevosto, V.: Verifying redundant-check based counter-
939 measures: a case study. In: *Proceedings of the 37th ACM/SIGAPP Symposium on*
940 *Applied Computing*. pp. 1849–1852 (2022)
- 941 69. Murdock, K., Oswald, D., Garcia, F.D., Van Bulck, J., Gruss, D., Piessens, F.:
942 Plundervolt: Software-based fault injection attacks against intel sgx. In: *2020 IEEE*
943 *Symposium on Security and Privacy (SP)*. pp. 1466–1482. IEEE (2020)
- 944 70. Mutlu, O., Kim, J.S.: Rowhammer: A retrospective. *IEEE Transactions on*
945 *Computer-Aided Design of Integrated Circuits and Systems* **39**(8), 1555–1571
946 (2019)
- 947 71. Niemetz, A., Preiner, M.: Bitwuzla at the SMT-COMP 2020. *CoRR*
948 **abs/2006.01621** (2020), <https://arxiv.org/abs/2006.01621>
- 949 72. Palazzi, L., Li, G., Fang, B., Pattabiraman, K.: A tale of two injectors: End-to-
950 end comparison of ir-level and assembly-level fault injection. In: *2019 IEEE 30th*
951 *International Symposium on Software Reliability Engineering (ISSRE)*. pp. 151–
952 162. IEEE (2019)
- 953 73. Papadakis, M., Malevris, N.: Automatic mutation test case generation via dynamic
954 symbolic execution. In: *2010 IEEE 21st International Symposium on Software Re-*
955 *liability Engineering*. pp. 121–130. IEEE (2010)
- 956 74. Pattabiraman, K., Nakka, N., Kalbarczyk, Z., Iyer, R.: Symplified: Symbolic
957 program-level fault injection and error detection framework. In: *2008 IEEE In-*
958 *ternational Conference on Dependable Systems and Networks With FTCS and*
959 *DCC (DSN)*. pp. 472–481. IEEE (2008)
- 960 75. Petrovic, G., Ivankovic, M., Kurtz, B., Ammann, P., Just, R.: An industrial appli-
961 cation of mutation testing: Lessons, challenges, and research directions. In: *2018*
962 *IEEE International Conference on Software Testing, Verification and Validation*
963 *Workshops (ICSTW)*. pp. 47–53. IEEE (2018)

- 964 76. Potet, M.L., Mounier, L., Puys, M., Dureuil, L.: Lazart: A symbolic approach
965 for evaluation the robustness of secured codes against control flow injections. In:
966 2014 IEEE Seventh International Conference on Software Testing, Verification and
967 Validation. pp. 213–222. IEEE (2014)
- 968 77. Preda, M.D., Giacobazzi, R., Debray, S., Coogan, K., Townsend, G.M.: Modelling
969 metamorphism by abstract interpretation. In: International Static Analysis Sym-
970 posium. pp. 218–235. Springer (2010)
- 971 78. Puys, M., Riviere, L., Bringer, J., Le, T.h.: High-level simulation for multiple fault
972 injection evaluation. In: Data Privacy Management, Autonomous Spontaneous Se-
973 curity, and Security Assurance, pp. 293–308. Springer (2014)
- 974 79. Rauzy, P., Guilley, S.: A formal proof of countermeasures against fault injection
975 attacks on crt-rsa. *Journal of Cryptographic Engineering* **4**(3), 173–185 (2014)
- 976 80. Recoules, F., Bardin, S., Bonichon, R., Lemerre, M., Mounier, L., Potet, M.L.:
977 Interface compliance of inline assembly: Automatically check, patch and refine. In:
978 2021 IEEE/ACM 43rd International Conference on Software Engineering (ICSE).
979 pp. 1236–1247. IEEE (2021)
- 980 81. Recoules, F., Bardin, S., Bonichon, R., Mounier, L., Potet, M.L.: Get rid of inline
981 assembly through verification-oriented lifting. In: 2019 34th IEEE/ACM Interna-
982 tional Conference on Automated Software Engineering (ASE). pp. 577–589. IEEE
983 (2019)
- 984 82. Richter-Brockmann, J., Sasdrich, P., Guneyasu, T.: Revisiting fault adversary
985 models—hardware faults in theory and practice. *IEEE Transactions on Comput-
986 ers* (2022)
- 987 83. Sen, K., Marinov, D., Agha, G.: Cute: A concolic unit testing engine for c. *ACM*
988 *SIGSOFT Software Engineering Notes* **30**(5), 263–272 (2005)
- 989 84. Shoshitaishvili, Y., Wang, R., Salls, C., Stephens, N., Polino, M., Dutcher, A.,
990 Groesen, J., Feng, S., Hauser, C., Kruegel, C., Vigna, G.: SoK: (State of) The Art
991 of War: Offensive Techniques in Binary Analysis. In: IEEE Symposium on Security
992 and Privacy (2016)
- 993 85. Skorobogatov, S.P., Anderson, R.J.: Optical fault induction attacks. In: Inter-
994 national workshop on cryptographic hardware and embedded systems. pp. 2–12.
995 Springer (2002)
- 996 86. Tang, A., Sethumadhavan, S., Stolfo, S.: {CLKSCREW}: Exposing the perils of
997 {Security-Oblivious} energy management. In: 26th USENIX Security Symposium
998 (USENIX Security 17). pp. 1057–1074 (2017)
- 999 87. Winter, S., Tretter, M., Sattler, B., Suri, N.: simfi: From single to simultaneous
1000 software fault injections. In: 2013 43rd Annual IEEE/IFIP International Confer-
1001 ence on Dependable Systems and Networks (DSN). pp. 1–12. IEEE (2013)
- 1002 88. <https://github.com/wookee-project>, accessed july 2021
- 1003 89. Zavalysyn, I., Given-Wilson, T., Legay, A., Sadre, R., Riviere, E.: Chaos duck: A
1004 tool for automatic iot software fault-tolerance analysis. In: 2021 40th International
1005 Symposium on Reliable Distributed Systems (SRDS). pp. 46–55. IEEE (2021)

1006 A Instrumentations of the motivating example

1007 We illustrate what would the instrumentation of the two main SWiFI techniques,
1008 mutant generation in Figure 7 and forking in Figure 8, look like for our motivat-
1009 ing example of Section 2. We also illustrate a forkless instrumentation in Figure
1010 9.

```
1 bool g_authenticated;  
2 int u1, u2, u3, u4, ref1, ref2, ref3, ref4;  
3  
4 void verifyPIN() {  
5     int res = non_det;  
6     res = res * (u1 == ref1);  
7     res = res * (u2 == ref2);  
8     res = res * (u3 == ref3);  
9     res = res * (u4 == ref4);  
10    g_authenticated = res;  
11 }  
12  
13 void main(int argc, char const *argv[]) {  
14     assert(u1!=ref1 && u2!=ref2 && u3!=ref3 && u4!=ref4);  
15     verifyPIN();  
16     assert(g_authenticated == true); /* Security oracle */  
17 }
```

```
1 void verifyPIN() {  
2     int res = 1;  
3     res = non_det;  
4     res = res * (u2 == ref2);  
5     res = res * (u3 == ref3);  
6     res = res * (u4 == ref4);  
7     g_authenticated = res;  
8 }
```

```
1 void verifyPIN() {  
2     int res = 1;  
3     res = res * (u1 == ref1);  
4     res = non_det;  
5     res = res * (u3 == ref3);  
6     res = res * (u4 == ref4);  
7     g_authenticated = res;  
8 }
```

Fig. 7: First three mutants of VerifyPIN function of the motivating example, for single fault only

```

1  bool g_authenticated;
2  int u1, u2, u3, u4, ref1, ref2, ref3, ref4;
3
4  /* booleans indicating if a fault happens */
5  bool b1, b2, b3, b4, b5, b6;
6  /* faulted values */
7  int non_det1, non_det2, non_det3, non_det4, non_det5,
8  non_det6;
9  /* Maximum number of faults in the fault model */
10 int max_f;
11
12 void verifyPIN() {
13     int res;
14     if (b1 <= max_f && b1)
15         then res = non_det1;
16     else res = 1;
17     if (b1 + b2 <= max_f && b2)
18         then res = non_det2
19     else res * (u1 == ref1);
20     if (b1 + b2 + b3 <= max_f && b3)
21         then res = non_det3 ;
22     else res * (u2 == ref2);
23     if (b1 + b2 + b3 + b4 <= max_f && b4)
24         then res = non_det4 ;
25     else res * (u3 == ref3);
26     if (b1 + b2 + b3 + b4 + b5 <= max_f && b5)
27         then res = non_det5 ;
28     else res * (u4 == ref4);
29     if (b1 + b2 + b3 + b4 + b5 + b6 <= max_f && b6)
30         then g_authenticated = non_det6 ;
31     else g_authenticated = res;
32 }
33
34 void main(int argc, char const *argv[]) {
35     assert(u1!=ref1 && u2!=ref2 && u3!=ref3 && u4!=ref4);
36     verifyPIN();
37     /* Security oracle */
38     assert(g_authenticated == true &&
39         (b1 + b2 + b3 + b4 + b5 + b6 <= max_f));
40 }

```

Fig. 8: Motivating example with forking instrumentation
(not exactly what happens in binsec, but would be my forking instrumentation)


```

1  bool g_authenticated;
2  int u1, u2, u3, u4, ref1, ref2, ref3, ref4;
3
4  /* booleans indicating if a fault happens */
5  bool b1, b2, b3, b4, b5, b6;
6  /* faulted values */
7  int non_det1, non_det2, non_det3, non_det4, non_det5,
8  non_det6;
9  /* Maximum number of faults in the fault model */
10 int max_f;
11
12 void verifyPIN() {
13     int res = ite b1 ? non_det1 : 1;
14     res = ite b2 ? non_det2 : res * (u1 == ref1);
15     res = ite b3 ? non_det3 : res * (u2 == ref2);
16     res = ite b4 ? non_det4 : res * (u3 == ref3);
17     res = ite b5 ? non_det5 : res * (u4 == ref4);
18     g_authenticated = ite b6 ? non_det6 : res;
19 }
20
21 void main(int argc, char const *argv[]) {
22     assert(u1!=ref1 && u2!=ref2 && u3!=ref3 && u4!=ref4);
23     verifyPIN();
24     /* Security oracle */
25     assert(g_authenticated == true &&
26           (b1 + b2 + b3 + b4 + b5 + b6 <= max_f));
27 }

```

Fig. 9: Motivating example with forkless instrumentation

1011 B Raw result tables

1012 We provide additional statistics in Tables 4 and 5, comparing our different algo-
1013 rithm implementations and the forking technique.

Table 4: Analysis time and explored paths comparison for 1 to 10 faults (RQ2-3)
Forking*: values for 3 faults and above are computed for incomplete runs (timeouts)
C: EDS+IOD

	Analysis time (s)							Explored paths						
	1f	2f	3f	4f	6f	8f	10f	1f	2f	3f	4f	6f	8f	10f
minimal value (0 means lower than 0.5s)														
FASE	0	0	0	0	0	0	0	6	6	6	6	6	6	6
EDS	0	0	0	0	0	0	0	6	6	6	6	6	6	6
IOD	0	0	0	0	0	0	0	6	6	6	6	6	6	6
C	0	0	0	0	0	0	0	6	6	6	6	6	6	6
Forking*	0	5	62	613	30k	86k	86k	36	516	5462	44k	1333k	1399k	1373k
average value														
FASE	3	6	10	13	16	18	18	18	57	109	164	236	261	269
EDS	2	5	8	13	16	19	18	18	57	109	164	236	261	269
IOD	1	3	5	8	10	10	10	18	57	109	164	236	261	269
C	1	3	6	8	10	10	10	18	57	109	164	236	261	269
Forking*	10	647	13k	40k	79k	86k	86k	215	11k	209k	764k	1788k	2087k	2086k
median value														
FASE	4	8	9	11	11	12	11	21	67	122	176	219	220	220
EDS	3	11	9	10	11	11	11	21	67	122	176	219	220	220
IOD	2	4	5	6	6	6	6	21	67	122	176	219	220	220
C	2	4	5	6	6	6	6	21	67	122	176	219	220	220
Forking*	14	1k	54k	86k	86k	86k	86k	234	14k	532k	1077k	1935k	2285k	2058k
maximal value														
FASE	24	156	733	2506	12k	25k	30k	42	337	1923	7351	46k	122k	172k
EDS	16	80	468	1839	11k	26k	31k	42	337	1923	7351	46k	122k	172k
IOD	7	44	250	943	5239	11k	12k	42	337	1923	7351	46k	122k	172k
C	8	50	287	1046	5639	11k	12k	42	337	1923	7351	46k	122k	172k
Forking	156	14k	86k	86k	86k	86k	86k	1k	96k	785k	2383k	2421k	3350k	3371k
number of timeouts (over 12 programs)														
FASE	0	0	0	0	0	0	0	0	0	0	0	0	0	0
EDS	0	0	0	0	0	0	0	0	0	0	0	0	0	0
IOD	0	0	0	0	0	0	0	0	0	0	0	0	0	0
C	0	0	0	0	0	0	0	0	0	0	0	0	0	0
Forking	0	0	6	9	11	12	12	0	0	6	9	11	12	12

Table 5: Average solving time per query and number of queries sent to the solver comparison for 1 to 10 faults (RQ3)

Forking*: values for 3 faults and above are computed for incomplete runs (timeouts)

Average solving time ($\times 10^{-3}$)								Queries sent to the solver							
	1f	2f	3f	4f	6f	8f	10f	1f	2f	3f	4f	6f	8f	10f	
minimal value															
FASE	4	6	7	8	7	9	8	11	11	11	11	11	11	11	
EDS	3	7	7	6	8	8	8	12	11	11	11	11	11	11	
IOD	3	5	5	6	5	6	5	12	12	12	12	12	12	12	
EDS+IOD	2	5	5	6	5	6	4	13	12	12	12	12	12	12	
Forking*	2	2	3	3	6	12	12	94	2093	22k	163k	3007k	3100k	2957k	
average value															
FASE	25	26	26	25	26	28	26	100	215	343	450	557	592	600	
EDS	18	24	23	26	25	28	25	88	203	335	458	579	611	617	
IOD	13	14	15	15	15	15	15	70	180	308	416	523	556	564	
EDS+IOD	12	14	15	15	15	15	14	81	205	341	452	544	572	579	
Forking*	11	15	17	17	18	18	19	844	42k	725k	2294k	4334k	4574k	4443k	
median value															
FASE	29	32	34	33	31	34	33	132	262	400	485	524	524	524	
EDS	18	23	29	34	34	37	33	128	251	392	508	547	540	540	
IOD	11	14	15	16	18	18	17	102	211	332	420	462	462	462	
EDS+IOD	12	14	16	16	18	17	15	117	244	386	476	482	478	478	
Forking*	14	18	20	20	19	18	19	1007	65k	1889k	3564k	4411k	4718k	4507k	
maximal value															
FASE	225	148	108	115	103	115	107	496	3470	15k	47k	203k	389k	456k	
EDS	138	197	141	107	98	115	102	394	3076	14k	47k	228k	456k	522k	
IOD	66	55	53	52	48	46	45	350	3001	15k	51k	256k	520k	617k	
EDS+IOD	65	65	53	51	49	47	45	372	3290	17k	58k	282k	572k	676k	
Forking*	25	30	31	27	28	27	29	6176	479k	3837k	9656k	6522k	7071k	6936k	
number of timeouts (over 12 programs)															
FASE	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
EDS	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
IOD	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
EDS+IOD	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
Forking	0	0	6	9	11	12	12	0	0	6	9	11	12	12	

1014 C Case-study, attacks details

1015 We provide additional details on attacks described in the WooKey bootloader
 1016 case study in Section 8. Figure 10 shows the relevant code extract.

1017 The relevant parts of the bootloader functions are presented in Figure 10.
 1018 Lacombe et al. found one attack in the *loader_exec_req_selectbank* function
 1019 at l.3 (A1) and one in the *loader_exec_req_flashlock* function at l.31 (A2).
 1020 They correspond to data corruption in branching conditions. Faulting the test l.3
 1021 when both firmware flip and flop are bootable results in the execution carrying
 1022 on to the test on l.7, which only checks if flop can be booted, assuming at

```

1 static loader_request_t loader_exec_req_selectbank (loaderstate_t
  nextstate){
2     // ...
3     if((flip_shared_vars.fw.bootable == FW BOOTABLE &&
4         flop_shared_vars.fw.bootable == FW BOOTABLE) &&
5         !(flip_shared_vars.fw.bootable != FW BOOTABLE ||
6         flop_shared_vars.fw.bootable != FW BOOTABLE)){
7         // ...
8     }
9     if(flop_shared_vars.fw.bootable == FW BOOTABLE
10    (CM1 - X) && flip_shared_vars.fw.bootable != FW BOOTABLE){
11        if(!(flop_shared_vars.fw.bootable == FW BOOTABLE
12            && flip_shared_vars.fw.bootable != FW BOOTABLE))
13            goto err;
14        ctx.boot_flop = sectrue ;
15        // ...
16    }
17    if(flip_shared_vars.fw.bootable == FW BOOTABLE
18    (CM3 - X) && flop_shared_vars.fw.bootable != FW BOOTABLE){
19        ctx.boot_flip = sectrue ;
20        // ...
21    }
22 }
23 static loader_request_t loader_exec_req_flashlock (loader_state_t
  nextstate){
24     // ...
25     if (ctx.dfu_mode == sectrue) {
26         (CMA) if (ctx.dfu_mode != sectrue)
27         (CMA) goto err;
28         // ...
29     }
30     else if (ctx.dfu_mode == secfalse) {
31         if (ctx.bootflip == sectrue) {
32             (CM4) if (ctx.bootflip != sectrue)
33             (CM4) goto err;
34             // ...
35             ctx.next_stage = (app_entry_t) (FW1_START);
36         }
37         // ...
38     }
39 }

```

Fig. 10: functions of WooKey’s bootloader, with [63] fixes and **our patch**

1023 least one firmware cannot. If both are bootable, and flop is the older one, the
 1024 attacker’s goal can be satisfied. The second attack described takes advantage
 1025 of the lack of countermeasure in *loader_exec_req_flashlock*, which computes
 1026 the pointer to the boot function of the chosen firmware. Inducing a fault that
 1027 inverts the test leads to the wrong pointer being selected. We are able to find
 1028 both attacks A1 and A2, linking faults back to their locations in the C code with
 1029 debug information.

1030 We find an additional attack, A3, faulting another part of the *loader_exec_*
 1031 *req_flashlock* function at line 25. By inverting that test, it is possible to select
 1032 the wrong firmware pointer to boot on, in particular, if we are not in a secure
 1033 Direct Firmware Update (DFU) mode, it is possible to boot as if we were.